CCIR · COMPUTE CREDIT INDEX RESEARCH

Governance and IOSCO Alignment Statement

Version 1.0.0 • March 2026

Administrator	CCIR Compute Credit Index Research
Contact	research@ccir.io
Website	ccir.io
Methodology	v1.0.0 (March 2026)
IOSCO Framework	Principles for Financial Benchmarks (IOSCO, July 2013)
IOSCO Alignment	Voluntary — maintained as precondition for credit document citation
Oversight Board	Targeted Q1 2027
Document Status	Active

This document is the formal governance statement for CCIR Compute Credit Index Research (“CCIR”), the administrator of the Compute Rate Index (“CRI”) family of GPU compute reference rates. It documents CCIR’s alignment with each of the 19 IOSCO Principles for Financial Benchmarks (July 2013) and is intended for review by structured finance counsel evaluating whether CRI rates may be cited in credit agreements, rating agency staff assessing benchmark credibility for criteria development, and lenders and credit analysts underwriting GPU-backed credit facilities.

IOSCO alignment is maintained voluntarily. CRI indices are not currently subject to mandatory benchmark regulation in any jurisdiction. Voluntary IOSCO alignment is maintained because it is the structural standard that structured finance counsel and rating agencies use to evaluate whether a benchmark is fit for citation in credit documents in a post-LIBOR regulatory environment where lender-controlled and marketplace-controlled rates are structurally disqualifying.

1. Administrator Profile and Structural Independence

CCIR Compute Credit Index Research is the sole administrator of the CRI index family. CCIR designs the index methodology, operates the data collection pipeline, executes index computation, and publishes all index values and accompanying disclosures. No element of benchmark determination is delegated to any third party.

CCIR holds no financial interest — equity, revenue share, referral fee, data licensing revenue, or other arrangement — in any GPU compute marketplace, cloud provider, or data center operator from which it collects data. This is structural independence: CCIR derives no revenue from any source it monitors and has no economic incentive to bias the indices toward any particular price level or market participant. This independence is the organizational precondition for CRI indices to function as citation-grade reference instruments in credit documents.

The post-LIBOR regulatory environment established a clear standard: a benchmark cited in a credit agreement must be produced by an administrator that is structurally independent from the market it measures. Lender-controlled metrics fail this standard because the lender has an economic interest in the benchmark level. Marketplace-produced rates fail this standard because the marketplace has a financial interest in the utilization and pricing of the assets it trades. CCIR's independence from all GPU marketplaces is the structural fact that resolves both disqualifying conditions.

CRI indices are sourced from six independently operated data providers across the spot index family (RunPod, Shadeform, Lambda Labs, TensorDock, Nebius, Verda) and three cluster providers (RunPod Cluster, Lambda Cluster, CoreWeave). Data collection is entirely algorithmic: prices are retrieved from public APIs and published rate card pages without human submitter involvement. No source communicates pricing to CCIR through any channel other than its public API or published rate card. No source is consulted, informed, or given advance notice of published index values.

2. IOSCO Principles — Summary Crosswalk

The table below summarizes CCIR's alignment status for each of the 19 IOSCO Principles for Financial Benchmarks. Three status designations are used. "Compliant" indicates full implementation as of March 2026. "Targeted" indicates a governance milestone with a specific date commitment. "Not Applicable" indicates a principle that is structurally inapplicable to CCIR's algorithmic collection model — in each such case, the inapplicability reflects a governance strength (absence of human submitters and the manipulation risks they introduce), not a gap. Detailed assessment for each principle follows in Section 3.

P#	IOSCO Principle	Status	CCIR Implementation Summary
P1	Overall Responsibility of the Administrator	**Compliant**	CCIR is sole administrator of all CRI indices, responsible for methodology design, pipeline operation, computation, and publication. Responsibility is not delegated to any third party.
P2	Oversight of Third Parties	**Compliant**	Azure Databricks infrastructure dependency is disclosed. Benchmark determination is not outsourced. Weekly audit_sources job validates all API data sources for authentication, required fields, and schema consistency.
P3	Conflicts of Interest for the Administrator	**Compliant**	CCIR holds no equity, revenue share, referral fee, or other financial interest in any GPU marketplace, cloud provider, or data center operator contributing data. Structural independence.
P4	Control Framework	**Compliant**	Six automated audit notebooks in /CCIR/Audit/: nightly audit_pipeline (post-publication validation), layer_reconciliation (Bronze→Silver→Gold row counts), delivery_manifest (SHA-256 tamper evidence), audit_digest_daily (HTML email digest); Sunday audit_sources (per-source API integrity), audit_digest_weekly (7-day summary email). FAIL conditions trigger job failure notifications. All results written to audit/ Delta tables in ADLS Gen2.

P#	IOSCO Principle	Status	CCIR Implementation Summary
P5	Internal Oversight	**Targeted — Q1 2027**	Formal Oversight Board formation targeted Q1 2027. Until formation, internal oversight is exercised through the six-notebook automated audit framework. Governance document will be revised upon Oversight Board formation.
P6	Benchmark Design	**Compliant**	Methodology fully documented with rationale for all design choices. Silver eligibility filters enforce economic comparability. Flat median across 7-day window. Per-GPU normalization. SXM and PCIe tracked separately. Cluster indices published alongside spot rates.
P7	Data Sufficiency	**Compliant**	Confidence flag (high/medium/low), source_count, and source_coverage (FULL/REDUCED/MINIMAL) published with every spot index value. All five spot indices are multi-source on the primary noninterruptible rate. Low-confidence values published with explicit disclosure, never suppressed.
P8	Hierarchy of Data Inputs	**Compliant**	Single input tier: algorithmically collected, publicly posted seller-ask prices passing Silver eligibility filters. No judgment-based or estimated inputs. Insufficient observations lower the confidence flag; no synthetic prices are constructed.
P9	Transparency of Benchmark Determinations	**Compliant**	Full methodology published at ccir.io. Source taxonomy, eligibility filters, and computation method fully disclosed. SHA-256 snapshot hashes link Gold values to specific collection events. Methodology version stamped on every row.
P10	Periodic Review	**Compliant**	Monthly model string review obligation documented in §7.1 of the methodology. Annual independent audit targeted for first full calendar year. Weekly audit_sources job provides continuous inter-period monitoring.
P11	Content of the Code of Conduct	**Not Applicable**	No human submitters. All data collected algorithmically from public APIs. The Silver filter layer encodes the eligibility decisions that would otherwise be subject to submitter judgment. Submitter manipulation and coordination risks are structurally absent.
P12	Submitter Code of Conduct	**Not Applicable**	No human submitters. Data collected algorithmically from public APIs and published rate cards. No individual exercises pricing discretion in the index determination process.
P13	Internal Controls Over Data Collection	**Compliant**	Bronze tables append-only, never modified after write. SHA-256 hashes computed per collection run and stamped on every Bronze row. Silver filters auditable against Bronze. Schema snapshots versioned in ADLS at utils/schema_snapshots/.
P14	Submitter Panels	**Not Applicable**	No submitter panel. Each data source contributes independently via its own API. No inter-source communication or price coordination is possible. Panel manipulation and anchoring risks are structurally absent.
P15	Complaints Procedures	**Compliant**	Complaints and methodology questions received at research@ccir.io. Written response within 10 business days. A formal complaints procedure has been published (March 2026) and is available at ccir.io. Upon Oversight Board formation (Q1 2027), the procedure will be updated to reflect escalation paths and the Board's role in reviewing disputed determinations.
P16	Audits	**Targeted — First Annual**	Annual independent third-party audit targeted for first full calendar year of publication. Six automated audit notebooks operational and writing to audit/ Delta tables: nightly

P#	IOSCO Principle	Status	CCIR Implementation Summary
			audit_pipeline, layer_reconciliation, delivery_manifest, audit_digest_daily; Sunday audit_sources, audit_digest_weekly.
P17	Audit Trail	**Compliant**	SHA-256 snapshot hashes on every Bronze row. Bronze append-only. Gold written via MERGE with idempotency guarantee. Delta Lake time travel supports point-in-time queries. audit/ Delta tables provide persistent audit records across all six audit notebooks. Methodology version stamped on every row.
P18	Cooperation with Regulatory Authorities	**Compliant**	IOSCO alignment is voluntary; CRI indices are not subject to mandatory benchmark regulation in any current jurisdiction. CCIR commits to cooperate fully with any regulatory authority asserting jurisdiction, including providing documentation and pipeline records on request.
P19	Outsourcing	**Compliant**	Data collection and computation infrastructure operates on Azure Databricks (Microsoft Azure) — disclosed. CCIR retains full and exclusive control of methodology design, eligibility filter definitions, computation logic, index publication, and governance. No element of benchmark determination is outsourced.

3. Detailed Principle Assessment

This section provides a detailed assessment of CCIR's alignment with each of the 19 IOSCO Principles, organized by IOSCO's four governance sections. Each principle is assessed against CCIR's current pipeline architecture, published methodology, and governance framework as of March 2026.

3.1 Section I — Governance (Principles 1–5)

Principle 1 Overall Responsibility of the Administrator — Compliant

CCIR Compute Credit Index Research is the sole administrator of all CRI indices. CCIR is solely responsible for methodology design, data collection pipeline operation, index computation, publication of index values and associated disclosures, governance framework maintenance, and periodic review. Responsibility is not delegated to any third party. There is no co-administrator arrangement.

The full methodology governing index construction is published in CCIR Methodology Statement v1.0.0 (March 2026), which documents all filter choices, computation methods, data source taxonomy, confidence flag definitions, and governance commitments with explicit rationale. Any change to a filter threshold, computation method, or source inclusion requires documentation of the rationale, a methodology version bump, and prospective application only. Historical values computed under a prior methodology version are not restated.

Principle 2 Oversight of Third Parties — Compliant

CCIR's data collection and computation infrastructure operates on Azure Databricks (Microsoft Azure). This infrastructure dependency is disclosed in the methodology and in this governance statement. Azure provides compute, storage (Azure Data Lake Storage Gen2), and key management (Azure Key Vault) infrastructure. Azure does not participate in any element of benchmark determination: methodology design, eligibility filter definitions, computation logic, and publication decisions are exclusively CCIR's.

Data source providers (RunPod, Shadeform, Lambda Labs, TensorDock, Nebius, Verda, CoreWeave, Vast.ai) are external API counterparties, not third-party delegates. CCIR monitors source API integrity through the weekly audit_sources job, which independently validates each source for: API authentication, required field presence, schema drift against versioned ADLS snapshots, new GPU model string detection, and source-specific structural integrity checks. Results are written to the audit/audit_log Delta table and require human review on the weekly cadence. The Sunday audit_digest_weekly email delivers a consolidated summary of the week's audit results including audit_sources findings to research@ccir.io.

Principle 3 Conflicts of Interest for the Administrator — Compliant

CCIR does not currently hold any equity interest, revenue share arrangement, referral fee, data licensing income, or other financial relationship with any GPU compute marketplace, cloud provider, or data center operator contributing data to the CRI pipeline. CCIR is not a market participant, broker, or platform operator. CCIR does not transact GPU compute on any basis. CCIR maintains internal policies to identify, disclose, and manage any potential future conflicts of interest as its activities develop.

This structural independence is the organizational fact that enables CRI indices to serve as citation-grade reference instruments in loan covenants and collateral appraisals. It is not a policy statement or a compliance declaration — it is a description of CCIR's organizational structure. A benchmark administrator whose revenue derives from the market it measures cannot serve as an independent reference instrument for credit documents. CCIR's revenue model does not include any component tied to GPU marketplace activity.

Principle 4 Control Framework — Compliant

CCIR operates an automated audit framework of six notebooks that runs independently from the data collection and index computation pipeline. The nightly audit_pipeline job validates published index CSV files end-to-end after each evening pipeline run, checking: column count integrity, today's row presence in each index file, methodology version consistency, index value sanity bounds (per-index floor and ceiling), observation count minimums, source count minimums, source coverage flags, observation inflation ratio, and confidence flag

plausibility. FAIL conditions raise exceptions that trigger Databricks job failure notifications. Warning conditions are logged but do not block publication.

The weekly audit_sources job independently validates each data source on a Sunday cadence, checking: API authentication, required field presence, schema drift against stored ADLS snapshots, and new GPU model string detection. Source-specific structural checks are tailored to the exact field names and API structures used in each Bronze collection notebook. Results from all audit jobs are written to a shared audit/audit_log Delta table, keyed on (audit_date, check_name, audit_source), providing a persistent and queryable audit record independent of job execution logs. Four additional notebooks complete the framework: layer_reconciliation (nightly — Bronze→Silver→Gold row count reconciliation via Delta transaction logs), delivery_manifest (nightly — SHA-256 checksums of all 11 published CSVs with two-pass tamper verification), audit_digest_daily (nightly — HTML email digest to research@ccir.io), and audit_digest_weekly (Sunday — 7-day summary email). All six audit notebooks write to the audit/ folder Delta tables in ADLS Gen2.

Principle 5 Internal Oversight — Targeted — Q1 2027

A formal Oversight Board is targeted for formation by Q1 2027. The Board is expected to include a majority of independent members with no financial interest in any GPU compute marketplace or data source contributing to the CRI pipeline. The Board's mandate will include review and challenge of proposed methodology changes, oversight of the annual audit process, and final determination on escalated complaints. The Q1 2027 target allows sufficient time-series accumulation across the full index family before the first formal board review, and provides adequate time to identify qualified independent members and formalize board charter, reporting obligations, and engagement terms. The governance document will be revised upon Oversight Board formation to reflect the board's composition, mandate, meeting cadence, and relationship to the audit framework.

Until Oversight Board formation, internal oversight is exercised by the CCIR administrator through the six-notebook automated audit framework described under Principle 4. This is a current gap relative to full IOSCO compliance, disclosed explicitly. Rating agency staff and structured finance counsel evaluating CCIR should weigh this gap against the compensating controls in place: six automated audit notebooks with persistent logging, append-only Bronze tables, SHA-256 audit trail, and full methodology transparency.

3.2 Section II — Benchmark Quality (Principles 6–10)

Principle 6 Benchmark Design — Compliant

CRI indices are designed to measure the seller-posted, noninterruptible, on-demand per-GPU-per-hour rental price for production-grade GPU compute across independent marketplaces and rate card providers. The design choices reflect the credit function the indices serve: noninterruptible pricing because an enterprise buyer absorbing redeployed collateral is not pricing interruptible capacity; per-GPU normalization because the GPU is the unit of collateral; SXM-only primary indices because SXM and PCIe are not price-equivalent and blending them would understate the premium SXM collateral commands.

Silver eligibility filters enforce economic comparability across all sources: full GPU allocation only, datacenter-grade reliability thresholds, single-GPU per-hour normalization, and physical machine deduplication on marketplace sources. The 7-day rolling window balances recency with observation stability — short enough to be responsive to market movements, long enough to smooth intraday fluctuations. The flat median aggregation is transparent, reproducible, and resistant to single-source dominance without opaque weighting. All design decisions are documented with rationale in the CCIR Benchmark Design Rationale and the Methodology Statement.

Cluster indices (CLS-MM and CLS-ENT) are published alongside spot rates to support spread analysis relevant to credit structuring. Generation spread monitoring (CRI-H100 relative to CRI-B200) provides a continuous signal of forward hardware obsolescence risk. These are additive disclosures; the core credit reference function is served by the five spot indices.

Principle 7 Data Sufficiency — Compliant

Every published index value carries a confidence flag — high, medium, or low — computed from observed source count and valid day count within the 7-day window. Thresholds are per-model, relative to SOURCE_EXPECTED (the methodology-locked expected source count for each gpu_model). High confidence requires source_count equal to SOURCE_EXPECTED on ≥ 5 valid days. Medium requires source_count ≥ ⌈SOURCE_EXPECTED / 2⌉ on ≥ 3 valid days. Low confidence values are published with full disclosure rather than suppressed, because a low-confidence value with 40 observations from a single source conveys different information than a low-confidence value with 3 observations, and both are more useful to a practitioner than silence.

All five spot indices carry multi-source noninterruptible coverage — six sources contribute observations across models (source mix varies by model; see the methodology Source Taxonomy table). No spot index carries structural low confidence on the primary published rate. CLS-ENT carries structural low confidence by design (CoreWeave is the sole enterprise source), disclosed in the bundle_note field of every ENT row and in the methodology. Source_count and total_observations are published daily alongside each index value.

Principle 8 Hierarchy of Data Inputs — Compliant

CRI indices use a single data input tier: algorithmically collected, publicly posted seller-ask prices that pass Silver eligibility filters. A hierarchy of data inputs is designed for benchmarks that may need to fall back from transaction data to indicative quotes to estimates when market liquidity is insufficient. CCIR's input is a single type: executable, machine-readable, publicly posted price observations. No judgment-based inputs, interpolated values, or estimated prices are used at any stage.

When Silver-eligible observations are insufficient for the standard computation, the confidence flag is lowered and the published value reflects whatever observations are available. No synthetic price is constructed to fill a gap. This is the methodologically conservative choice for a credit reference rate: a low-confidence published value with full disclosure gives a lender the information needed to assess suitability; a synthetic price obscures the gap.

Principle 9 Transparency of Benchmark Determinations — Compliant

The full CCIR methodology is published at ccir.io and in the CCIR Methodology Statement. Source taxonomy, eligibility filters for each source, computation method (flat median, 7-day window), confidence flag definitions, cluster index construction, and governance framework are all fully disclosed. Every published index value is accompanied by its confidence flag, source_count, total_observations, and as_of_date.

SHA-256 snapshot hashes are computed from the raw API response for each collection run and stamped on every Bronze row. These hashes link every published Gold index value to the specific collection events that generated the underlying Silver observations. Bronze tables are append-only and never modified after write, so the audit trail from a published index value back to the raw API response is preserved and queryable. Methodology version is stamped on every Bronze, Silver, and Gold row, creating a complete provenance record. A code-to-methodology traceability matrix (audit/traceability Delta table) maps each methodology section to the specific notebook path, cell reference, and code anchor that implements it, providing a direct link between the published methodology document and the running production code.

Principle 10 Periodic Review — Compliant

CCIR's methodology review obligations are documented in §7.1 of the Methodology Statement. GPU API naming conventions change without notice, and a drop in observation count for any model is the primary indicator of a gpu_name string change at the source. Source-specific model variant registries are reviewed monthly against Bronze gpu_model_raw values. Any mismatch is corrected and documented in the methodology changelog.

Annual independent audit is targeted for the first full calendar year of publication. The weekly audit_sources job provides continuous inter-period monitoring. Any methodology change triggered by a periodic review requires documentation of the rationale, a version bump, and prospective application only. The governance document will be updated annually or upon any material governance change.

3.3 Section III — Quality of the Methodology (Principles 11–14)

Principles 11, 12, and 14 address the design, content, and oversight of submitter codes of conduct and submitter panels. These principles are structurally inapplicable to CCIR. CCIR collects all data algorithmically from public APIs and published rate card pages. There are no human submitters who exercise pricing discretion in the index determination process. This is documented explicitly rather than omitted, because the absence of human submitters is a governance strength — it eliminates the submitter manipulation, anchoring, and coordination risks that motivated these IOSCO principles.

Principle 11 Content of the Code of Conduct — Not Applicable

This principle requires benchmarks that use human submitters to define and publish a code of conduct governing submitter behavior, conflicts, and escalation procedures. CCIR does not use human submitters. All price observations are collected algorithmically from public APIs and published rate card sources. No individual exercises pricing discretion, judgment, or authority in the index determination process.

The functional analog in CCIR's design is the Silver eligibility filter layer. Silver filters encode the comparability and quality decisions that would otherwise require submitter discipline: full GPU allocation requirements, datacenter-grade reliability thresholds, non-interruptible pricing only, and physical machine deduplication. These rules are fixed in code, applied uniformly to all sources, and auditable against the underlying Bronze data. There is no submitter who could deviate from them.

Principle 12 Submitter Code of Conduct — Not Applicable

This principle requires that submitters to a benchmark operate under a published code of conduct binding their submission behavior. CCIR has no submitters. Data is collected by CCIR's automated pipeline directly from source APIs; the data sources are not submitters in the IOSCO sense, because they do not submit price observations to CCIR for inclusion in an index. CCIR retrieves publicly posted prices from their APIs. The sources have no knowledge of how their pricing is used in the CRI computation.

Principle 13 Internal Controls Over Data Collection — Compliant

Bronze tables are append-only Delta tables that are never modified after the initial write. All Silver eligibility filters are applied exclusively at the Silver layer, meaning every filter decision is auditable against the unmodified raw Bronze data. SHA-256 snapshot hashes are computed from the raw API response for each collection run and stamped on every Bronze row, creating a tamper-evident link between each observation and the specific API call that produced it.

Schema snapshots for each data source are stored in ADLS at utils/schema_snapshots/{source}_schema.json and updated when a schema change is detected and confirmed as intentional. This creates a versioned record of each source API's response structure over time, making schema drift detectable on the weekly audit cadence. The weekly audit_sources job validates required field presence against these snapshots and flags new or removed fields for human review.

Principle 14 Submitter Panels — Not Applicable

This principle addresses the governance of submitter panels — the groups of financial institutions that submit data to benchmarks like bank rate panels. CCIR does not operate a submitter panel. Each data source contributes independently via its own public API. No inter-source communication exists, no source is aware of other sources' pricing, and no coordination mechanism of any kind is present. Panel manipulation and anchoring risks, which motivated this IOSCO principle, are structurally absent from CCIR's architecture.

3.4 Section IV — Accountability (Principles 15–19)

Principle 15 Complaints Procedures — Compliant

Methodology questions, data quality concerns, source inclusion disputes, and formal complaints are received at research@ccir.io. CCIR commits to acknowledge all complaints within 2 business days and to provide a written determination within 10 business days of receipt. Complaints and resolutions are logged internally.

A formal complaints procedure has been published (March 2026) and is available at ccir.io alongside this governance statement. The procedure specifies complaint scope, review timelines, interim disclosure obligations, and corrective action standards. Upon Oversight Board formation (Q1 2027), the procedure will be updated to reflect escalation paths and the Board's role in reviewing disputed determinations.

Principle 16 Audits — Targeted — First Annual

CCIR intends to obtain an annual independent audit conducted under a recognized assurance standard, targeting completion in the first full calendar year of publication. The audit will cover methodology adherence, data collection integrity, computation accuracy, and governance framework compliance. The audit firm will be selected on the basis of documented independence criteria; the firm's name and engagement scope will be disclosed upon appointment.

Internal automated audits are fully operational in the current period. The nightly audit_pipeline job validates the full collection-to-publication chain after each evening pipeline run. The weekly audit_sources job independently validates each data source. Nightly layer_reconciliation and delivery_manifest notebooks provide pipeline row-count integrity and SHA-256 tamper evidence respectively. Nightly audit_digest_daily and Sunday audit_digest_weekly deliver consolidated HTML audit summaries to research@ccir.io. All six audit notebooks write results to the audit/ Delta tables in ADLS Gen2. These internal controls are a compensating control during the period before the first external audit.

Principle 17 Audit Trail — Compliant

CCIR maintains a comprehensive, persistent audit trail from raw API collection through published index values. SHA-256 snapshot hashes are computed from the raw API response for each collection run and stamped on every Bronze row. Because Bronze tables are append-only, the raw data underlying any published value is permanently preserved and queryable. Every Silver and Gold row carries the snapshot hash of the Bronze row that contributed to it, maintaining traceability across pipeline layers.

Gold index values are written via MERGE keyed on (as_of_date, gpu_model, price_type_canonical) for spot indices and (as_of_date, gpu_model, cls_tier) for cluster indices, preventing row proliferation on re-run and making each daily computation idempotent. Methodology version is stamped on every Bronze, Silver, and Gold row, so any published value can be matched to the specific methodology version that governed its computation. Delta Lake time travel supports point-in-time queries of pipeline state at any historical date. The audit/ Delta tables in ADLS Gen2 record all automated audit results independently of job execution logs. A code-to-methodology traceability matrix (audit/traceability Delta table) maps each methodology section to the specific notebook path, cell reference, and code anchor that implements it.

Principle 18 Cooperation with Regulatory Authorities — Compliant

CRI indices are not currently subject to mandatory benchmark regulation in any jurisdiction. The EU Benchmark Regulation (BMR), the UK BMR, and relevant US regulatory frameworks do not currently extend mandatory requirements to GPU compute benchmarks. CCIR's IOSCO alignment is maintained voluntarily, as a structural precondition for credit document citation, not in response to a regulatory obligation.

CCIR commits to cooperate fully with any regulatory authority that asserts jurisdiction over its benchmark activities. This commitment includes providing access to methodology documentation, pipeline architecture documentation, audit logs, Bronze data, and any other records requested by a regulator with jurisdiction. CCIR will update this governance statement and its operational procedures promptly upon any regulatory development that affects the CRI index family.

Principle 19 Outsourcing — Compliant

CCIR's data collection and computation infrastructure operates on Azure Databricks, Microsoft Azure's managed analytics platform. This infrastructure dependency is disclosed in the methodology and in this governance statement. Azure Databricks provides managed compute, Delta Lake storage, and job scheduling infrastructure. CCIR retains full and exclusive control of all elements that constitute benchmark determination: methodology design, Silver eligibility filter definitions, computation logic, source selection and exclusion decisions, confidence flag computation, and index publication.

No element of index determination is outsourced. The pipeline code, filter definitions, and computation notebooks are authored, maintained, and controlled exclusively by CCIR. Azure executes CCIR-authored code on CCIR-controlled compute resources; Azure does not make any determination about which observations are eligible, what the computation method is, or what values are published. If CCIR were to migrate to a different compute infrastructure, the methodology and pipeline output would be identical.

4. Governance Milestones

The table below summarizes the current status of CCIR's governance framework components and targeted milestones. "Operational" indicates full implementation as of March 2026. "Targeted" indicates a committed milestone with a specific date. "In progress" indicates an active initiative without a fixed completion date.

Governance Item	Timing	Status	Detail
Automated audit framework	Nightly + weekly	**Operational**	Six notebooks in /CCIR/Audit/: nightly audit_pipeline, layer_reconciliation, delivery_manifest, audit_digest_daily; Sunday audit_sources, audit_digest_weekly. Results written to audit/ Delta tables in ADLS Gen2.
Methodology v1.0.0 publication	March 2026	**Complete**	ccir.io
Data Source Quality Policy publication	March 2026	**Complete**	ccir.io. IOSCO P14 Not Applicable functional replacement.
Monthly model string review	Monthly	**Operational**	§7.1 of methodology
Annual independent audit	First full calendar year	**Targeted**	Third-party audit firm TBD
Oversight Board formation	Q1 2027	**Targeted**	Composition and charter TBD
Formal complaints procedure	March 2026	**Complete**	Published March 2026. Available at ccir.io alongside this governance statement. Escalation path to Oversight Board will be added upon formation (Q1 2027).

5. Regulatory Disclaimer

CRI indices are not regulated financial benchmarks under any mandatory regulatory framework as of March 2026. The EU Benchmark Regulation (2016/1011), the UK Benchmarks Regulation, and relevant provisions of US financial regulation do not currently extend mandatory requirements to GPU compute reference rates. CCIR's alignment with the IOSCO Principles for Financial Benchmarks is voluntary and is maintained as a precondition for credit document citation, not in response to any regulatory mandate.

Nothing in this document constitutes legal advice, financial advice, or a representation that CRI indices are suitable for any particular use in any particular credit agreement or jurisdiction. Users of CRI indices in credit documents should obtain independent legal advice on the applicability of benchmark regulation and the enforceability of benchmark references in their specific agreements and jurisdictions.

CCIR acts solely as a benchmark administrator. CCIR does not owe fiduciary duties to any party referencing CRI rates in credit documents, and nothing in this governance statement or in any published CRI rate creates a fiduciary relationship between CCIR and any lender, borrower, rating agency, or other party.

To the fullest extent permitted by applicable law, CCIR shall not be liable for any direct, indirect, incidental, or consequential loss or damage arising from the use of, reliance on, or inability to use any published CRI rate, including any loss arising from a Force Majeure Cessation, a methodology change, or an error in any published index value. Users of CRI rates in credit documents assume all risk associated with such use.

6. Document Control

Version	1.0.0
Date	March 2026
Supersedes	N/A — initial issuance
Next Review	March 2027, or upon material governance change
Administrator	CCIR Compute Credit Index Research
Contact	research@ccir.io

CCIR Compute Credit Index Research · research@ccir.io · ccir.io · This document describes CCIR's governance framework and voluntary IOSCO alignment as of March 2026.